

SQL Queries (iii)

- Sets in SQL
- Bags in SQL
- The **IN** Operator
- The **EXISTS** Operator
- Quantifiers
- Union, Intersection, Difference
- Division
- Selection with Aggregation

❖ Sets in SQL

The relational model is set-based

Set literals are written as $(expr_1, expr_2, \dots)$ (each $expr_i$ yields an atomic value)

SQL query results are (more or less) sets of tuples or atomic values

Examples:

```
-- set literals  
(1,2,3)      ('a','b','c','d')  
-- set of atomic values  
(select salary from Employees)  
-- set of tuple values  
(select id, name from Employees)
```

SQL provides a variety of set-based operators: \in , \cup , \cap , $-$, $/$, \exists , \forall , ...

❖ Bags in SQL

SQL query results are actually bags (multisets), allowing duplicates, e.g.

```
select age from Students;  
-- yields (18,18,18,...19,19,19,19,...20,20,20,...)
```

Can convert bag to set (eliminate duplicates) using **DISTINCT**, e.g.

```
select distinct age from Students;
```

SQL set operations **UNION**, **INTERSECT**, **EXCEPT** ...

- yield sets by default (i.e. eliminate duplicates)
- can produce bags with keyword **ALL** (e.g. **UNION ALL**)

```
(1,2,3) UNION (2,3,4) yields (1,2,3,4)  
(1,2,3) UNION ALL (2,3,4) yields (1,2,3,2,3,4)
```

❖ The IN Operator

Tests whether a specified tuple is contained in a relation
(i.e. $t \in R$)

tuple **IN** *relation* is true iff the tuple is contained in the relation.

Conversely for *tuple* **NOT IN** *relation*.

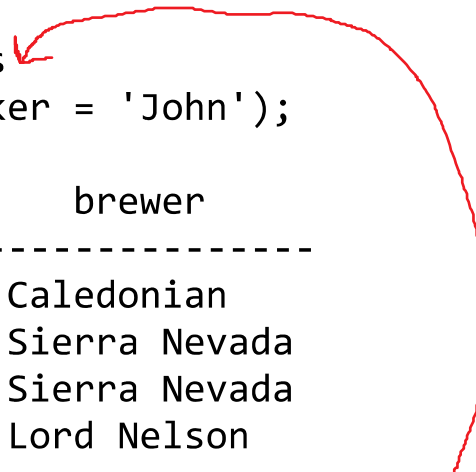
Syntax:

```
SELECT *  
FROM   R  
WHERE  R.a IN (SELECT x FROM S WHERE Cond)  
        -- assume multiple results
```

❖ The IN Operator (cont)

Example: Find the name and brewer of beers that John likes.

```
SELECT name, brewer
FROM   Beers
WHERE  name IN
      (SELECT beer
       FROM   Likes
       WHERE  drinker = 'John');
```



name	brewer
80/-	Caledonian
Bigfoot Barley Wine	Sierra Nevada
Pale Ale	Sierra Nevada
Three Sheets	Lord Nelson

Subquery = "What are the names of the beers that John likes?"

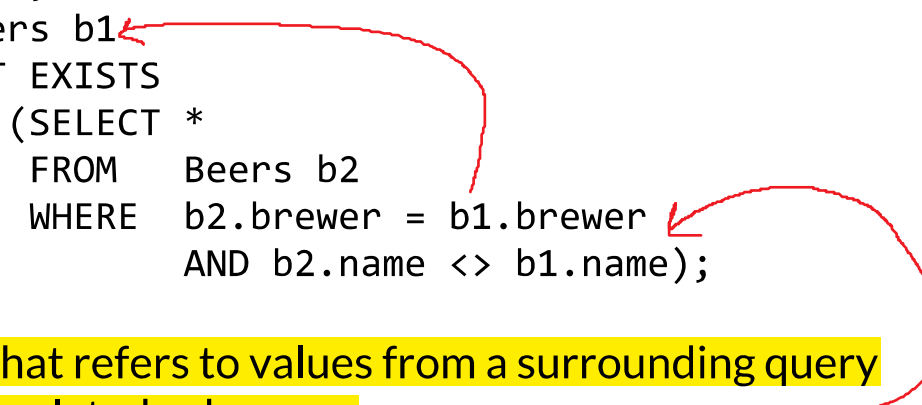
(This and subsequent beer queries use an older smaller version of the Beer database)

❖ The EXISTS Operator

EXISTS(*relation*) is true iff the relation is non-empty.

Example: Find the beers that are the unique beer by their manufacturer.

```
SELECT name, brewer
FROM   Beers b1
WHERE  NOT EXISTS
      (SELECT *
       FROM   Beers b2
       WHERE  b2.brewer = b1.brewer
              AND b2.name <> b1.name);
```



A subquery that refers to values from a surrounding query is called a correlated subquery.

❖ Quantifiers

ANY and **ALL** behave as **existential** and **universal** quantifiers respectively.

Example: Find the beers sold for the highest price.

```
SELECT beer
FROM   Sells
WHERE  price >=
        ALL(SELECT price FROM sells);
```

Beware: in common use, "any" and "all" are often synonyms.

E.g. "I'm better than any of you" vs. "I'm better than all of you".

❖ Union, Intersection, Difference

SQL implements the standard set operations

$R1 \text{ UNION } R2$ set of tuples in either $R1$ or $R2$

$R1 \text{ INTERSECT } R2$ set of tuples in both $R1$ and $R2$


$R1 \text{ EXCEPT } R2$ set of tuples in $R1$ but not $R2$

$R1$ and $R2$ must be union-compatible (i.e. same schema)

Union and intersection semantics are straightforward.

❖ Union, Intersection, Difference (cont)

Example: Find the drinkers and beers such that the drinker likes the beer and frequents a bar that sells it.



```
(SELECT drinker, beer FROM Likes)
INTERSECT
(SELECT drinker, beer
 FROM Sells natural join Frequents);
```

drinker		beer
Adam		New
John		Three Sheets
Justin		Victoria Bitter

❖ Union, Intersection, Difference (cont)

Set difference is implemented by **EXCEPT**

R

A	B
1	'a'
2	'b'
3	'a'

S

A	B
1	'a'
1	'b'
2	'a'

R except S

A	B
2	'b'
3	'a'

S except R

A	B
1	'b'
2	'a'

Semantics of set difference: $R \text{ except } S = \{ x \in R, \text{ where } x \notin S \}$

for each entry in R, if it is in S then don't include it in the result

COMP3311 20T3 ♦ SQL Queries (iii) ♦ [9/13]

R except S:

1:a is in S so don't include
2:b is not in S, include
3:a is not in S, include

therefore result of R except S is 2:b, 3:a

S except R:

1:a is in R, so don't include
1:b is not in R, include
2:a is not in R, include

therefore result of S except R is 1:b, 2:a

❖ Division

Division aims to find values in one table that occur in conjunction with all values in another table:

R		S		R / S	
A	B	B		A	
1	'a'	'a'		1	
2	'b'	'b'		2	
3	'a'				
1	'b'				
2	'a'				

Arises in queries like "Find Xs that are related to all Ys / every Y"

COMP3311 20T3 ♦ SQL Queries (iii) ♦ [10/13]

in table R, only 1 and 2 have an entry for both 1 and 2. 3 only has an entry with a, hence it is missing b and so wont show up in the division.

Division is a bit of a weird name. I like to think of it as "divided out" i.e. each entry of B can be assigned to the same entry in A in the other table i.e. i can give a to 1 and i can give b to 1, hence i can give all of the entries out and the division is ==1. For 3 on the other hand, i can only dish out the a, the b cannot be divided out, hence it is only 0.5 divided out, and hence not a proper division.

❖ Division (cont)

Not all SQL implementations provide a division operator

But can be achieved by combination of existing operations

Example: Find bars that each sell all of the beers Justin likes.

```

SELECT DISTINCT a.bar  select all the bars (no duplicates) where
FROM    Sells a
WHERE   NOT EXISTS (
            (SELECT beer FROM Likes
              WHERE drinker = 'Justin')  all the beers Justin likes
          EXCEPT
            (SELECT beer FROM Sells b
              WHERE bar = a.bar)  all the beers that are sold at the current bar
        );

```

i.e. is (all my liked beers) - (this bars beers) = 0?

COMP3311 20T3 ♦ SQL Queries (iii) ♦ [11/13]

remember the except is: for each beer in R, if it is also in S, remove from the resulting set. Hence if there is anything in the resulting set, then this bar doesn't sell all the beers he likes, hence the not exists returns false and this bar won't be included.

Another way of looking at it is the resulting set starts as being all the beers justin likes i.e. R. You then remove a beer from the resulting set if it is also in S. Hence there are at minimum 0 beers (all get removed), or the result set doesn't change from its initial state (i.e. it remains = R).

How is this a division? Well the bars are being divided by all beers justin likes i.e. only the bars that have a mapping to EVERY beer that justin likes will remain in the result set.

❖ Selection with Aggregation

Selection clauses can contain aggregation operations.

Example: What is the average price of New?

```
SELECT AVG(price)
FROM   Sells
WHERE  beer = 'New';
```

avg

2.38749998807907

- the bag semantics of SQL gives the correct result here
- the price for New in all hotels will be included, even if two hotels sell it at the same price
- if we used set semantics, we'd get the average of all the **different** prices for New.

❖ Selection with Aggregation (cont)

If we want set semantics, can force using **DISTINCT**.

Example: How many different bars sell beer?

```
SELECT COUNT(DISTINCT bar)
FROM   Sells;
```

```
count
-----
      6
```

Without **DISTINCT**, counts number of entries in the **Sells** table.

Aggregation operators on numbers: **SUM**, **AVG**, **MIN**, **MAX**,

Produced: 4 Oct 2020