

## PLpgSQL (iii)

---

- PLpgSQL Functions (recap)
- Query results in PLpgSQL
- Dynamically Generated Queries
- Functions vs Views

## ❖ PLpgSQL Functions (recap)

---

Defining PLpgSQL functions:

```
CREATE OR REPLACE
    funcName(param1, param2, ....)
    RETURNS rettype
AS $$
DECLARE
    variable declarations
BEGIN
    code for function
END;
$$ LANGUAGE plpgsql;
```

## ❖ Query results in PLpgSQL

Can evaluate a query and iterate through its results

- one tuple at a time, using a **for ... loop**

```
declare
    tup Type;
begin
    for tup in Query
    loop
        Statements;
    end loop;
end;
```

Type of **tup** variable must match type of *Query* results

If declared as **record**, will automatically match *Query* results type

## ❖ Query results in PLpgSQL (cont)

**Example:** count the number of Employees earning more than min.salary

```
create or replace function
    well_paid(_minsal integer) returns integer
as $$
declare
    nemps integer := 0;
    tuple record; -- could also be tuple Employees;
begin
    for tuple in
        select * from Employees where salary > _minsal
    loop
        nemps := nemps + 1;
    end loop;
    return nemps;
end;
$$ language plpgsql;
```

## ❖ Query results in PLpgSQL (cont)

Alternative to the above (but less efficient):

```
create or replace function
    well_paid(_minsal integer) returns integer
as $$
declare
    nemps integer := 0;
    tuple record;
begin
    for tuple in
        select name,salary from Employees
    loop
        if (tuple.salary > _minsal) then
            nemps := nemps + 1;
        end if;
    end loop;
    return nemps;
end;
$$ language plpgsql;
```

COMP3311 20T3 ♦ PLpgSQL (iii) ♦ [4/11]

If you can do the logic in the where condition of the query, then do that, because it is more efficient. It's not hard to see that we will return all employee records here, and so we are needlessly iterating over ones that don't factor into the answer, and we have to do the if check every loop because of it.

## ❖ Query results in PLpgSQL (cont)

And the example could be done more simply (and efficiently) as:

```
create or replace function
    well_paid(_minsal integer) returns integer
as $$
declare
    nemps integer;
begin
    select count(*) into nemps
    from Employees where salary > _minsal
    return nemps;
end;
$$ language plpgsql;
```

COMP3311 20T3 ♦ PLpgSQL (iii) ♦ [5/11]

ayyy, I was thinking about using count!

now we don't need a loop at all!

I'm not sure if this is true, but perhaps you can think of the SQL queries as being  $O(1)$ , and so putting the count and where condition into the query, instead of using a loop and if check, speeds it up by making it  $O(1)$  overall and not  $O(n) \cdot O(1) = O(n)$  (with the loop and if check).

## ❖ Dynamically Generated Queries

**EXECUTE** takes a string and executes it as an SQL query.

Examples:

```
-- constant string
execute 'select * from Employees';
-- concatenation of constant strings
execute 'select * from ' || 'Employees';
-- using a name of e.g. table or attribute
execute 'select * from ' || quote_ident($1);
-- using a value generated in the program
execute 'delete from Accounts' ||
       ' where holder=' || quote_literal($1);
```

Can be used in any context where an SQL query is expected

This mechanism allows us to construct queries "on the fly".

## ❖ Dynamically Generated Queries (cont)

Example: a wrapper for updating a single text field

```
create or replace function
    set(_table text, _attr text, _val text) returns void
as $$
declare
    query text;
begin
    query := 'update ' || quote_ident(_table);
    query := query || ' SET ' || quote_ident(_attr);
    query := query || ' = ' || quote_literal(_val);
    execute query;
end; $$ language plpgsql;
```

which could be used as e.g.

```
select set('branches', 'assets', '0.00');
```



## ❖ Dynamically Generated Queries (cont)

### One limitation of EXECUTE:

- cannot use **select into** inside dynamic queries

### Needs to be expressed instead as:

```
declare tuple R%rowtype; n int;
execute 'select * from R where id=' || n into tuple;
-- or
declare x int; y int; z text;
execute 'select a,b,c from R where id=' || n into x,y,z;
```

### Notes:

- if query returns multiple tuples, first one is stored
- if query returns zero tuples, all nulls are stored

## ❖ Functions vs Views

A difference between views and functions returning a **SETOF**:

- **CREATE VIEW** produces a "virtual" table definition
- **SETOF** functions require an existing tuple type

In examples above, we used existing **Employees** tuple type.

In general, you need to define the tuple return type via

```
create type TupleType as ( attr1 type1, ... attrn typen );
```

Other major differences between **setof** functions and views ...

- functions have parameters; views don't (although **where** might help)
- functions are "run-time" objects; views are interpolated into queries

## ❖ Functions vs Views (cont)

Another example of function returning **setof** tuples...

```
create type EmpInfo as (name text, pay integer);

create or replace function
    richEmps(_minsal integer) returns setof EmpInfo
as $$
declare
    emp record;    info EmpInfo;
begin
    for emp in
        select * from Employees where salary > _minsal
    loop
        info.name := emp.name;
        info.pay := emp.salary;
        return next info;
    end loop;
end; $$ language plpgsql;
```

## ❖ Functions vs Views (cont)

Using the function ...

```
select * from richEmps(100000);
```

versus a view

```
create or replace view richEmps(name,pay) as view is hard-coded  
select name, salary from Employees where salary > 100000;
```

```
select * from richEmps; -- but no scope for different salary
```

versus an SQL function

```
create or replace function  
    richEmps(_minsal integer) returns setof EmpInfo  
as $$  
select name, salary from Employees where salary > _minsal;  
$$ language sql;
```

Produced: 10 Oct 2020